

ARTIFICIAL NEURAL NETWORKS

FINAL PROJECT REPORT

About Dataset:

This dataset is originally from the National Institute of Diabetes and Digestive and Kidney Diseases. The objective of the dataset is to diagnostically predict whether a patient has diabetes, based on certain diagnostic measurements included in the dataset. Several constraints were placed on the selection of these instances from a larger database. All patients here are females at least 21 years old of Pima Indian heritage.

Predict whether the patients in the dataset have diabetes or not?

The dataset contains total of 768 records and each record contains 9 features and all the columns are listed below:

Features:

Pregnancies - Number of times pregnant.

Glucose - Plasma glucose concentration 2 hours in an oral glucose tolerance test.

Blood Pressure - Diastolic blood pressure (mm Hg).

Skin Thickness - Triceps skin fold thickness (mm).

Insulin - 2-Hour serum insulin (mu U/ml).

BMI - Body mass index (weight in kg/ (height in m) ^2).

Diabetes Pedigree Function - Diabetes pedigree function.

Age - Age (years).

Outcome - Class variable (0 or 1) 268 of 768 are 1, the others are 0.

Describing Dataset:

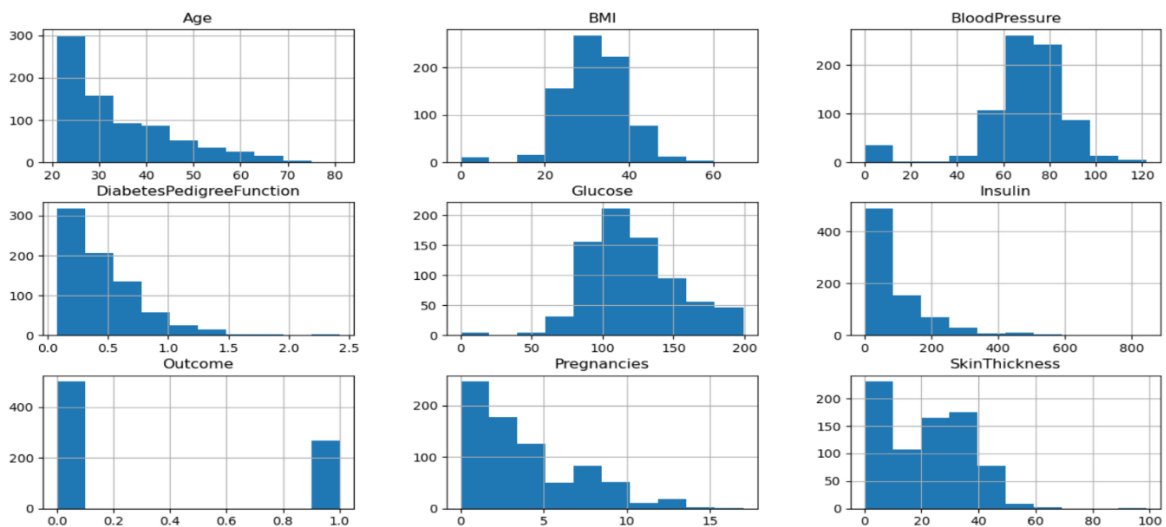
Before going any further, we should first check how the data has been distributed and get a little more information on the data. This can be done using describe function in pandas. From the image below we can say that the average age of the patient is 33 years old and minimum age is 21 years and Blood Pressure is also an interesting feature where average blood pressure stood at 69.10 and where 75 percent of the people are in ideal blood pressure range of 60mmHg - 80 mmHg. This gave more insights on the data.

```
In [9]: df.describe()
```

```
Out[9]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.471876	33.240885	0.348958
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	21.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.243750	24.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.372500	29.000000	0.000000
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.626250	41.000000	1.000000
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000	1.000000

Also check the data has been distributed across each feature (here we left out the outcome feature)



Data Preprocessing:

Why is it necessary to preprocess the data?

Because there is always chance of data containing null values or any other missing values or you can remove unwanted features and data that cannot be used, and it will lead bad results. So, it is always good practice to preprocess and clean the data before feeding it to the Machine Learning Algorithm.

For this dataset I have preprocessed the data and I have replaced the values that 0 with average values of that feature.

```

-----
Before preprocessing
Number of rows with 0 values for each variable
Pregnancies: 111
Glucose: 5
BloodPressure: 35
SkinThickness: 227
Insulin: 374
BMI: 11
DiabetesPedigreeFunction: 0
Age: 0
Outcome: 500
-----

After preprocessing
Number of rows with 0 values for each variable
Pregnancies: 111
Glucose: 0
BloodPressure: 0
SkinThickness: 0
Insulin: 0
BMI: 0
DiabetesPedigreeFunction: 0
Age: 0
Outcome: 500
-----

```

Splitting the Data:

Why should we split the data?

If we don't split the data the problem is that we are training the model on the dataset and again we are testing the model with same dataset so in this case we will get good results but if we test the same model on the new data it may not perform best as it has performed on the current dataset. So, it is always good practice to split the data into training and test set.

Here for this dataset I am splitting it into training and test as show in the image below:

```

In [15]: # Split the data into a training and testing set
X = df.loc[:, df.columns != 'Outcome']
y = df.loc[:, 'Outcome']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

```

After splitting the data, I am feeding the data to a neural network as defined below:

```
In [16]: # Build neural network in Keras
model = Sequential()
model.add(Dense(32, activation='relu', input_dim=8))
model.add(Dense(16, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
model.fit(X_train, y_train, epochs=200, verbose=False)
```

As you can see from the code I have used a simple neural network to train the model it has 2 hidden layers one with 32 nodes and activation function is relu and other hidden layer with 16 nodes activation function as relu and the final node is the output layer where I am using sigmoid to get the binary output. I am also compiling the model using adam as my optimizer with the default learning rate 0.001.

How to Check if our model has performed better?

Because it is a classification problem, we can use accuracy and f1 score to determine how well our model has done.

Accuracy Results:


Training Accuracy: 89.90%

Testing Accuracy: 81.17%

F1Score:

0.6947368421052632

ROC Curve:



Conclusion:

From the results of accuracy and f1score I can say that my model has performed better in giving results 70 percent of the times but we can improve the model performance and results if we train the by feeding it more data.